

Ordering Read and Write Operations on a Multicore CPU



José Paumard

PHD, JAVA CHAMPION, JAVA ROCK STAR

@JosePaumard <http://blog.paumard.org/>



Agenda



We saw what **synchronization** is

This module is about **visibility**

Bound to the way **multicore CPUs** work

What is the **happens before** link

What does **volatile** mean

The impact of **false sharing** on code



Synchronization and Visibility



Synchronization

Synchronization protects a block of code

Guarantees this code is executed by one thread at a time

Prevents race condition



Race Condition

```
public void consume() {  
    synchronized(lock) {  
        while (isEmpty(buffer)) {}  
        buffer[--count] = 0;  
    }  
}
```

```
public void produce() {  
    synchronized(lock) {  
        while (isFull(buffer)) {}  
        buffer[count++] = 1;  
    }  
}
```



Race Condition

count



```
public void consume() {  
    synchronized(lock) {  
        while (isEmpty(buffer)) {}  
        buffer[--count] = 0;  
    }  
}
```

```
public void produce() {  
    synchronized(lock) {  
        while (isFull(buffer)) {}  
        buffer[count++] = 1;  
    }  
}
```



Race Condition

count



```
public void consume() {  
    synchronized(lock) {  
        while (isEmpty(buffer)) {}  
        buffer[--count] = 0;  
    }  
}
```

- 1) reads count from memory
- 2) decrement it
- 3) writes count back to memory



Race Condition

count



- 1) reads count from memory
- 2) increment it
- 3) writes count back to memory

```
public void produce() {  
    synchronized(lock) {  
        while (isFull(buffer)) {}  
        buffer[count++] = 1;  
    }  
}
```



Memory Access

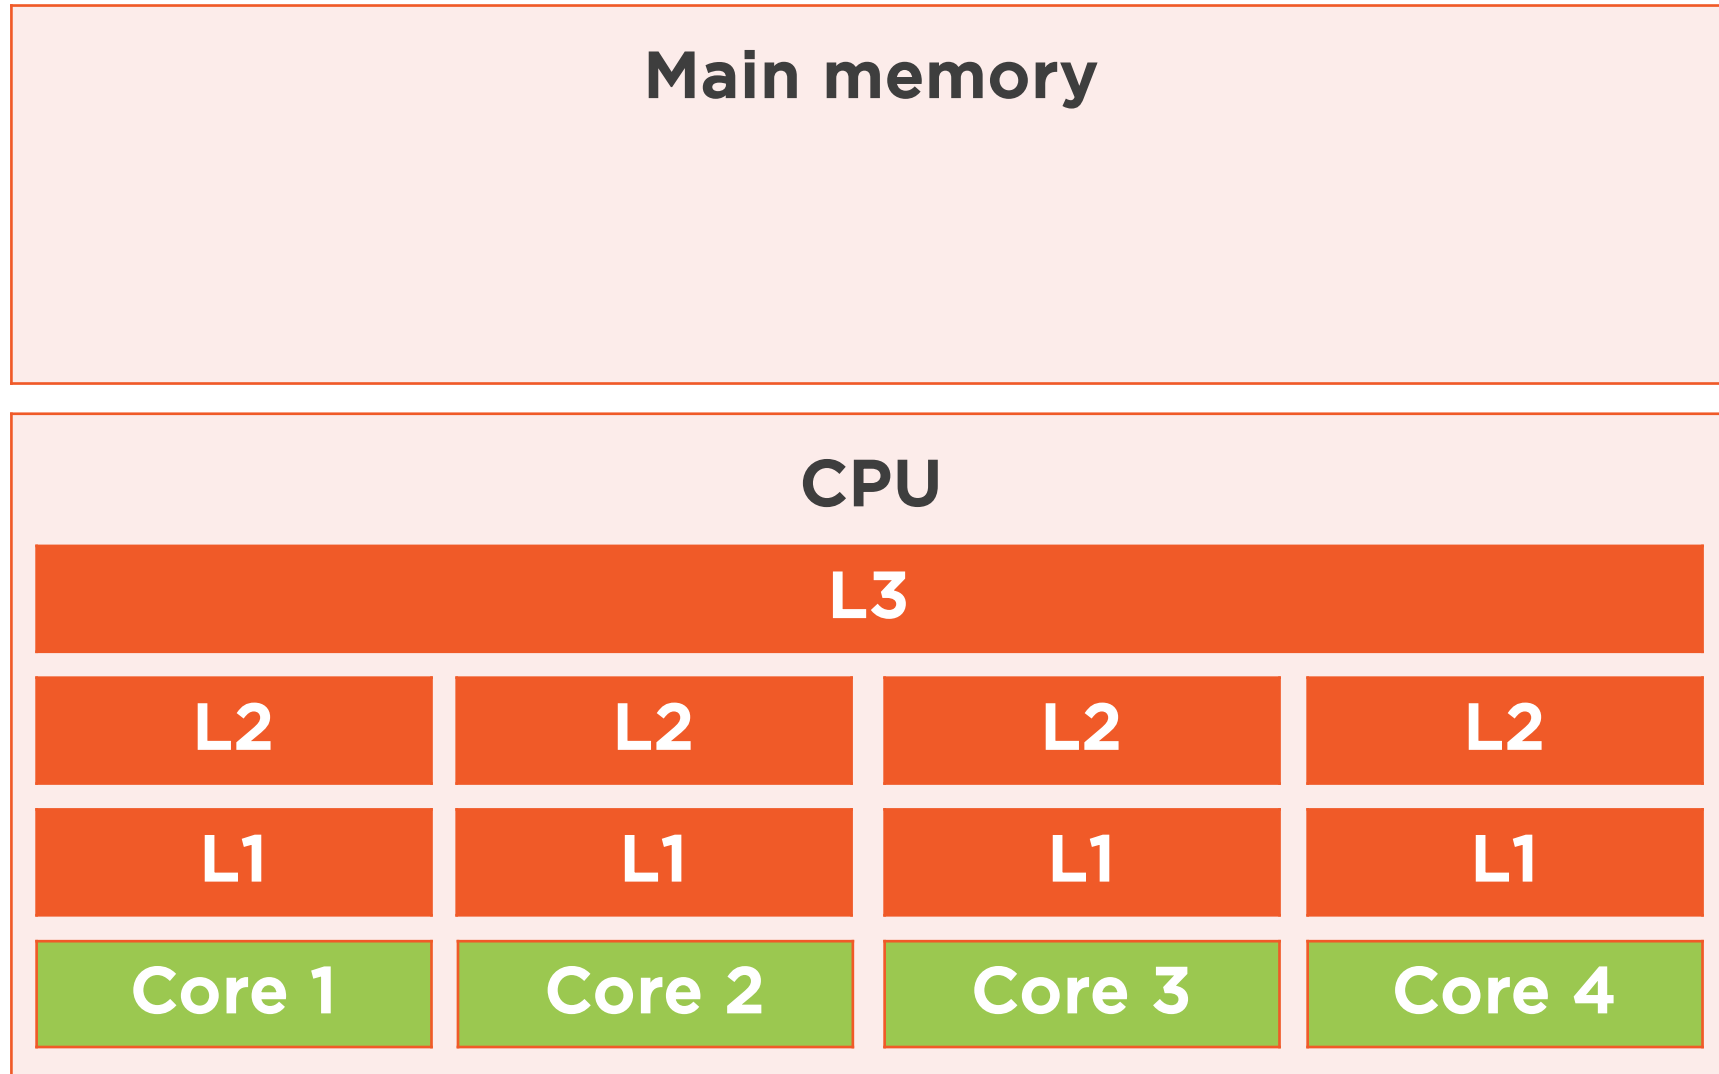
20 years ago, when CPU had no cache, this code was working fine

But nowadays, things do not work like that any more!

A CPU does not read a variable from the main memory, but from a cache



CPU Architecture



Why Has It Been Made Like That?

Because access to caches is much faster
than access to main memory

Access to the main memory: ~100ns

Access to the L2 cache: 7ns

Access to the L1 cache: 0.5ns



There Are Tradeoffs

Size of the main memory: several GB

Size of the L2 cache: 256kB

Size of the L1 cache: 32kB



CPU Architecture

Main memory

count


0

Core 1 needs count

Core 1



CPU Architecture





Core 1 needs count

1) The variable is copied in L1



CPU Architecture

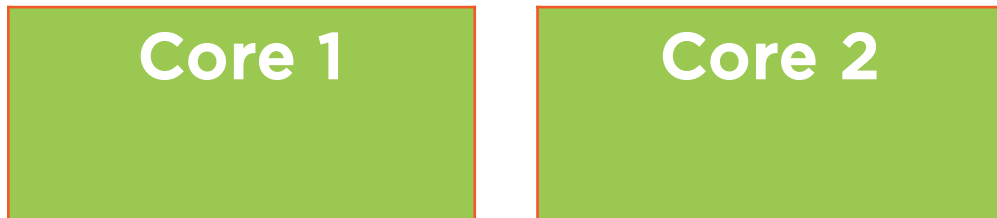



Core 1 needs count

- 1) The variable is copied in L1
- 2) Core 1 can modify it



CPU Architecture





Core 1 needs count

1) The variable is copied in L1

2) Core 1 can modify it

3) Core 2 also needs count

CPU Architecture



Core 1 needs count

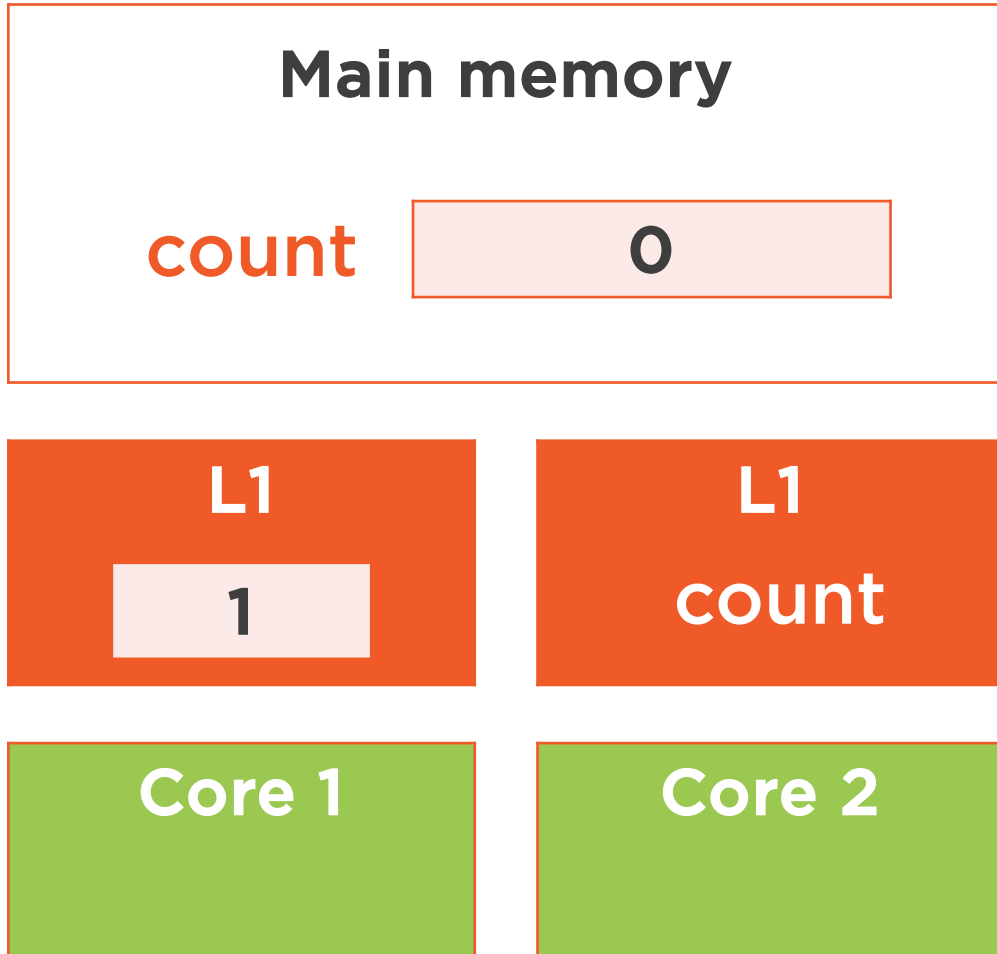
1) The variable is copied in L1

2) Core 1 can modify it

3) Core 2 also needs count

4) It should get the value 1, not 0

CPU Architecture



Core 1 needs count

1) The variable is copied in L1

2) Core 1 can modify it

3) Core 2 also needs count

4) It should get the value 1, not 0

This is visibility!

Visibility

A variable is said **visible**

If the writes made on it are **visible**

All the **synchronized writes** are **visible**



“Happens Before” Link



Happens Before

There are multiple references to *happens-before* in the Javadoc

Memory consistency effects: As with other concurrent collections, actions in a thread prior to placing an object into a `BlockingQueue` *happen-before* actions subsequent to the access or removal of that element from the `BlockingQueue` in another thread.

This interface is a member of the Java Collections Framework.

Since:

1.5

Memory consistency effects: Actions in a thread prior to calling `await()` *happen-before* actions that are part of the barrier action, which in turn *happen-before* actions following a successful return from the corresponding `await()` in other threads.

Since:

1.5

Memory consistency effects: Until the count reaches zero, actions in a thread prior to calling `countDown()` *happen-before* actions following a successful return from a corresponding `await()` in another thread.

Since:

1.5



The Java Memory Model

Multicore CPU brings new problems

Read and writes can really happen at the
same time

A given variable can be stored in more than
one place

Visibility means “a read should return the
value set by the last write”

What does last mean in a multicore world?



The Java Memory Model

We need a timeline to put read and write operations on



The Java Memory Model

Thread T_1

$x = 1$

1) T_1 writes 1 to x



The Java Memory Model

Thread T_1

$x = 1$

1) T_1 writes 1 to x

2) T_2 reads x and copy it to r

Thread T_2

$r = x$



The Java Memory Model

Thread T_1

$x = 1$

- 1) T_1 writes 1 to x
- 2) T_2 reads x and copy it to r
- 3) What is the value of r ?

Thread T_2

$r = x$



The Java Memory Model

Thread T_1

$x = 1$

Thread T_2

$r = x$

- 1) T_1 writes 1 to x
- 2) T_2 reads x and copy it to r
- 3) What is the value of r ?

The Java Memory Model fixes the answer to this question



The Java Memory Model

Thread T_1

$x = 1$


If there is no “happens before” link between the two operations, the value of r is unknown

Thread T_2

$r = x$



The Java Memory Model



If there is no “happens before” link between the two operations, the value of r is unknown

If there is a “happens before” link between the two operations, the value of r is 1

How Can We
Set up a
“Happens
Before” Link?

**There is no such keyword in the Java
language...**



Happens before link

A “happens before” link exists between all synchronized or volatile write operations and all synchronized or volatile read operations that follow



Happens Before Examples

```
int index;  
  
void increment() {  
    index++;  
}  
  
void print() {  
    System.out.println(index);  
}
```

Two operations:

- 1) a write operation
- 2) a read operation



Happens Before Examples

```
int index;  
  
void increment() {  
    index++;  
}  
  
void print() {  
    System.out.println(index);  
}
```

Two operations:

- 1) a write operation
- 2) a read operation

What does this code print in multithread?



Happens Before Examples

```
int index;  
  
void increment() {  
    index++;  
}  
  
void print() {  
    System.out.println(index);  
}
```

Two operations:

- 1) a write operation
- 2) a read operation

What does this code print in multithread?

No synchronization, no volatility → impossible to say



Happens Before Examples

```
int index;  
  
void synchronized increment() {  
    index++;  
}  
  
void synchronized print() {  
    System.out.println(index);  
}
```

Two operations:

- 1) a write operation
- 2) a read operation

What does this code print in multithread?

Synchronization → the correct value is always printed



Happens Before Examples

```
volatile int index;  
  
void increment() {  
    index++;  
}  
  
void print() {  
    System.out.println(index);  
}
```

Two operations:

- 1) a write operation
- 2) a read operation

What does this code print in multithread?

The variable is volatile → the correct value is always printed



A More Complex Example

```
int x, y, r1, r2;  
Object lock = new Object();
```

```
void firstMethod() {  
    x = 1;  
    synchronized(lock) {  
        y = 1;  
    }  
}
```

```
void secondMethod() {  
    synchronized(lock) {  
        r1 = y;  
    }  
    r2 = x;  
}
```

firstMethod() is writing x and y
secondMethod() is reading them

They are executed in
threads T_1 and T_2

Question: what is the value of r2?



A More Complex Example

```
int x, y, r1, r2;  
Object lock = new Object();
```

```
void firstMethod() {  
    x = 1;  
    synchronized(lock) {  
        y = 1;  
    }  
}
```

```
void secondMethod() {  
    synchronized(lock) {  
        r1 = y;  
    }  
    r2 = x;  
}
```

Due to the way the code is written, there is a *happens-before* link between:

- $x = 1$ and $y = 1$
- $r1 = y$ and $r2 = x$



A More Complex Example

```
int x, y, r1, r2;  
Object lock = new Object();
```

```
void firstMethod() {  
    x = 1;  
    synchronized(lock) {  
        y = 1;  
    }  
}
```

```
void secondMethod() {  
    synchronized(lock) {  
        r1 = y;  
    }  
    r2 = x;  
}
```

If T_1 is the first to enter the synchronized block, then the execution is in this order:

- $x = 1$

- $y = 1$

- $r1 = y$

- $r2 = x$

Happens-before link between a synchronized write and a synchronized read

The value of $r2$ is 1



A More Complex Example

```
int x, y, r1, r2;  
Object lock = new Object();
```

```
void firstMethod() {  
    x = 1;  
    synchronized(lock) {  
        y = 1;  
    }  
}
```

```
void secondMethod() {  
    synchronized(lock) {  
        r1 = y;  
    }  
    r2 = x;  
}
```

If T_2 is the first to enter the synchronized block, then the execution is in this order:

- $r1 = y$

- $r2 = x$ or $x = 1$?

- $y = 1$

No happens-before link
between $r2 = x$ and $x = 1$

The value of $r2$ may be 0 or 1



Synchronization

Guarantees the exclusive execution of a block of code



Synchronization

Guarantees the exclusive execution of a block of code

Visibility

Guarantees the consistency of the variables



All shared variables should be accessed in a synchronized or a volatile way



False Sharing



What Is “False Sharing”?

False sharing happens because of the way the CPU caches work

It is a side effect, that can have a tremendous effect on performance



What Is “False Sharing”?

The cache is organized in lines of data

Each line can hold 8 longs (64 bytes)

When a visible variable is modified in an L1 cache, all the line is marked “dirty” for the other caches

A read on a dirty line triggers a refresh on this line

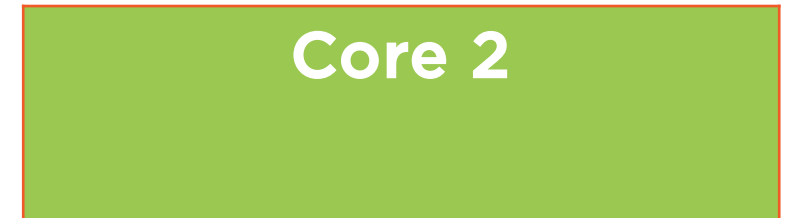
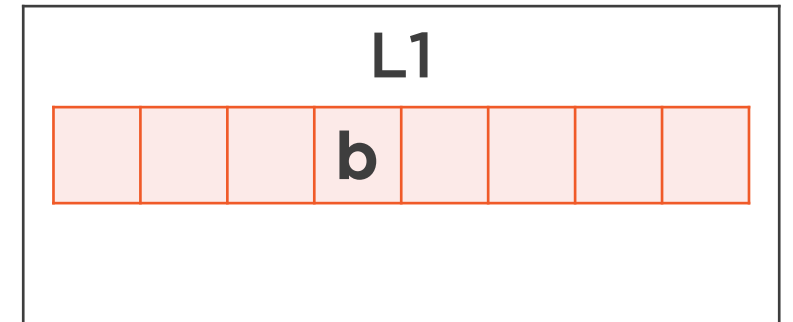
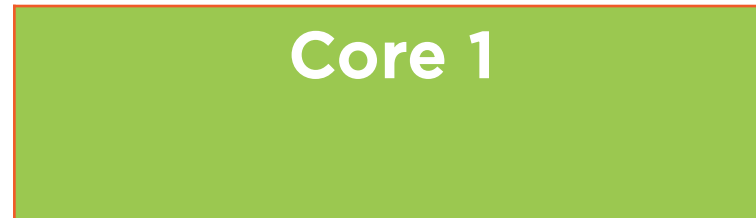
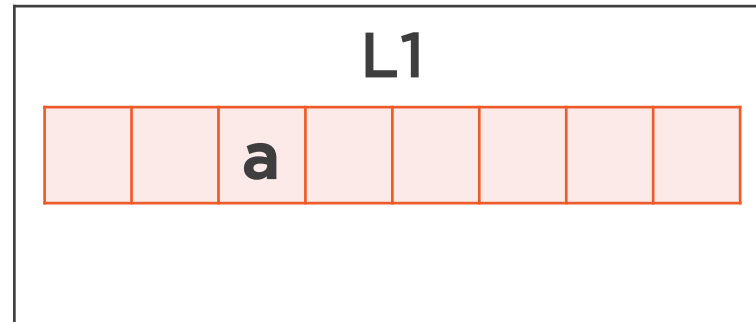


From Main Memory to L1 Cache

```
volatile long a, b;
```

```
void firstMethod() {  
    a++;  
}
```

```
void secondMethod() {  
    b++;  
}
```







From Main Memory to L1 Cache

```
volatile long a, b;
```

```
void firstMethod() {  
    a++;  
}
```

```
void secondMethod() {  
    b++;  
}
```


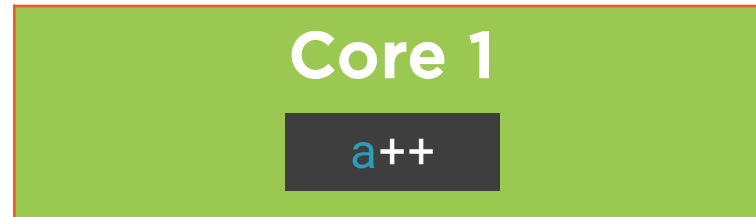
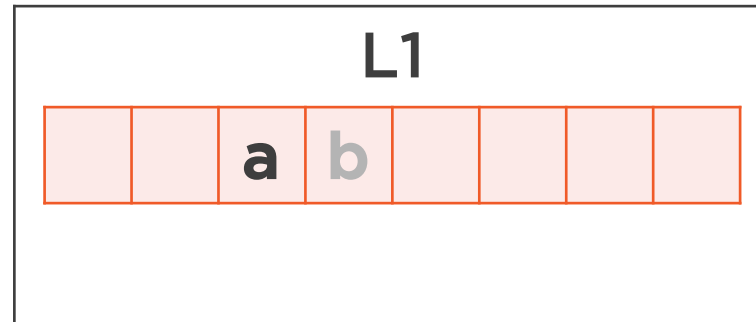


From Main Memory to L1 Cache

```
volatile long a, b;
```

```
void firstMethod() {  
    a++;  
}
```

```
void secondMethod() {  
    b++;  
}
```




From Main Memory to L1 Cache

```
volatile long a, b;
```

```
void firstMethod() {  
    a++;  
}
```

```
void secondMethod() {  
    b++;  
}
```



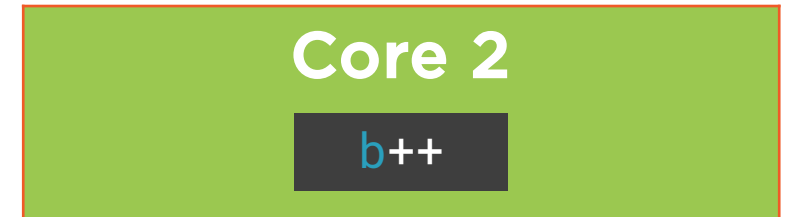
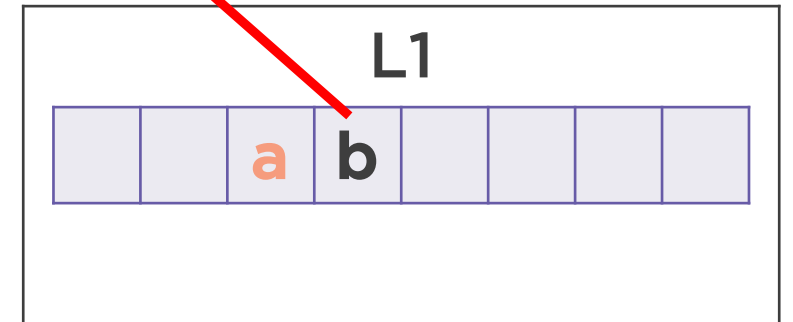
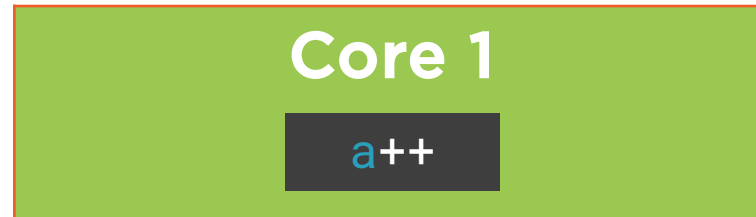
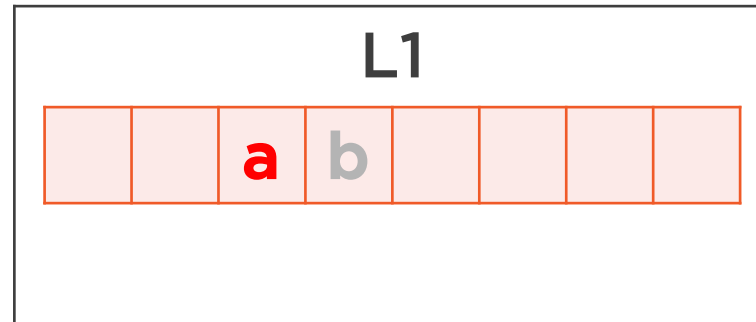
From Main Memory to L1 Cache

```
volatile long a, b;
```

```
void firstMethod() {  
    a++;  
}
```

```
void secondMethod() {  
    b++;  
}
```

Cache miss,
go to main memory

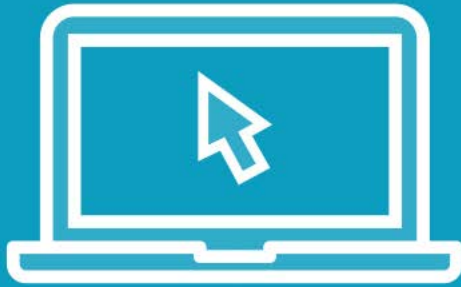


False Sharing

False sharing happens in an invisible way
Hard to predict, hits your performance
There are workarounds though...



Demo



Let us see some code!

An example of false sharing and the influence of variable padding



Wrapup



What did we learn?

The importance of visibility in multicore CPU

Two fundamental notions in concurrent applications: synchronization and volatility

How the way caching works in CPU can generate false sharing

