

# Using Locks and Semaphores for the Producer / Consumer Pattern

---



**José Paumard**

PHD, JAVA CHAMPION, JAVA ROCK STAR

@JosePaumard <https://github.com/JosePaumard>



# Agenda



About **synchronization**

**Synchronized & volatile:** **intrinsic locking**

**Explicit locking with Lock**

**Wait / notify with Lock and Condition**

**Semaphores**



# Intrinsic / Explicit Locking

---



```
public class Person {  
    private final Object key = new Object();  
  
    public String init() {  
        synchronized(key) {  
            // do some stuff  
        }  
    }  
}
```

This code prevents more than one thread to execute the synchronized block at the same time



```
public class Person {  
    private final Object key = new Object();  
  
    public String init() {  
        synchronized(key) {  
            // do some stuff  
        }  
    }  
}
```

What happens if several threads are trying to execute the `init()` block?

One of them will be allowed in the block, the others will have to wait for their turn



```
public class Person {  
    private final Object key = new Object();  
  
    public String init() {  
        synchronized(key) {  
            // do some stuff  
        }  
    }  
}
```

What happens if a thread is “blocked” inside the block?

All the other threads are also blocked

There is no way to release them...



# The Lock Pattern Brings a Richer API to Handle This Case



```
Object key = new Object();
```

```
synchronized(key) {  
    // do some stuff  
}
```

Instead of writing this code





```
Lock lock = new ReentrantLock();  
try {  
    lock.lock();  
    // do some stuff  
} finally {  
    lock.unlock();  
}
```

We write this code

Lock is an interface, implemented by ReentrantLock

It offers the same guarantees (exclusion, read & write ordering)

... and more functionalities!



## Let Us Analyze This Pattern

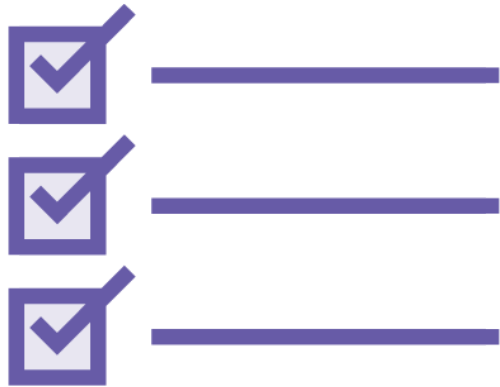
Instead of passing an instance of **Object**

To guard **several blocks** of code

We create a **lock object**

With **methods** on it





Interruptible lock acquisition

Timed lock acquisition

Fair lock acquisition



```
Lock lock = new ReentrantLock();  
try {  
    lock.lock();  
    // do some stuff  
} finally {  
    lock.unlock();  
}
```

This is the basic pattern



```
Lock lock = new ReentrantLock();  
try {  
    lock.lockInterruptibly();  
    // do some stuff  
} finally {  
    lock.unlock();  
}
```

This is the interruptible pattern

The thread will wait until it can enter the guarded block of code

But another thread can interrupt it by calling its `interrupt()` method

This can be costly, or hard to achieve though...





Interruptible lock acquisition

Timed lock acquisition

Fair lock acquisition



```
Lock lock = new ReentrantLock();  
if (lock.tryLock()) {  
    try {  
        // guarded block of code  
    } finally {  
        lock.unlock();  
    }  
} else { ... }
```

This is a timed lock acquisition

If a thread is already executing the guarded block of code

Then tryLock() returns false, immediately



```
Lock lock = new ReentrantLock();  
if (lock.tryLock(1, TimeUnit.SECONDS)) {  
    try {  
        // guarded block of code  
    } finally {  
        lock.unlock();  
    }  
} else { ... }
```

This is a timed lock acquisition

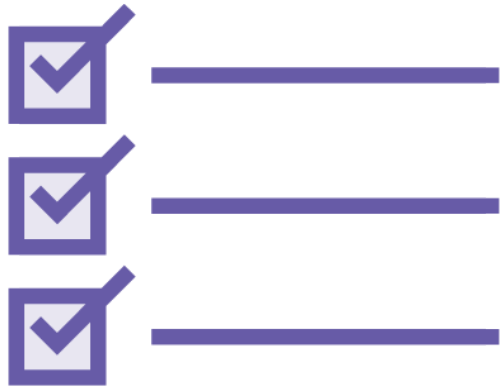
If a thread is already executing the guarded block of code

Then tryLock() returns false, immediately

One can also pass a time out as a parameter





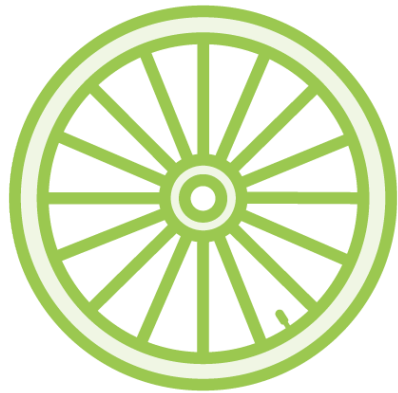


Interruptible lock acquisition

Timed lock acquisition

Fair lock acquisition





When several threads are waiting for a lock

Whether it is an intrinsic or explicit lock...

The first one to enter the guarded block of code is chosen randomly

Fairness means that the first to enter the wait line is the first to enter the block of code



```
Lock lock = new ReentrantLock();  
try {  
    lock.lock();  
    // guarded block of code  
} finally {  
    lock.unlock();  
}
```

A ReentrantLock built in the normal way is non-fair



```
Lock lock = new ReentrantLock(true); // fair
try {
    lock.lock();
    // guarded block of code
} finally {
    lock.unlock();
}
```

A ReentrantLock built in the normal way is non-fair

And one can also pass a boolean to it

True: this lock is fair, false: this lock is non-fair

A fair lock is costly...



# Using Lock Gives Wiggle Room

A lock can be interrupted: possible but hard, costly

It can be blocked for a set amount of time

It can be fair, letting in the threads on a first come first served basis



# Implementing the Producer / Consumer Pattern

---





Intrinsic locks: the producer / consumer pattern is implemented with wait / notify

Obviously, wait / notify cannot work with explicit locking

So we need another pattern!



```
Object lock = new Object();
```

```
class Producer {  
    public void produce() {  
        synchronized(lock) {  
            while (isFull(buffer))  
                lock.wait();  
            buffer[count++] = 1;  
            lock.notifyAll();  
        }  
    }  
}
```





```
Object lock = new Object();
```

```
class Producer {  
    public void produce() {  
        synchronized(lock) {  
            while (isFull(buffer))  
                lock.wait();  
            buffer[count++] = 1;  
            lock.notifyAll();  
        }  
    }  
}
```

```
class Consumer {  
    public void consume() {  
        synchronized(lock) {  
            while (isEmpty(buffer))  
                lock.wait();  
            buffer[--count] = 0;  
            lock.notifyAll();  
        }  
    }  
}
```



```
Object lock = new Object();
```

```
class Producer {  
    public void produce() {  
        synchronized(lock) {  
            while (isFull(buffer))  
                lock.wait();  
            buffer[count++] = 1;  
            lock.notifyAll();  
        }  
    }  
}
```

```
class Consumer {  
    public void consume() {  
        synchronized(lock) {  
            while (isEmpty(buffer))  
                lock.wait();  
            buffer[--count] = 0;  
            lock.notifyAll();  
        }  
    }  
}
```



```
Object lock = new Object();
```

```
class Producer {  
    public void produce() {  
        synchronized(lock) {  
            while (isFull(buffer))  
                lock.wait();  
            buffer[count++] = 1;  
            lock.notifyAll();  
        }  
    }  
}
```

```
class Consumer {  
    public void consume() {  
        synchronized(lock) {  
            while (isEmpty(buffer))  
                lock.wait();  
            buffer[--count] = 0;  
            lock.notifyAll();  
        }  
    }  
}
```



```
Lock lock = new ReentrantLock();
```

```
class Producer {  
    public void produce() {  
        try {  
            lock.lock();  
            while (isFull(buffer))  
                // wait  
            buffer[count++] = 1;  
            // notify  
        } finally {  
            lock.unlock();  
        }  
    }  
}
```



```
Lock lock = new ReentrantLock();
```

```
class Producer {  
    public void produce() {  
        try {  
            lock.lock();  
            while (isFull(buffer))  
                // wait  
            buffer[count++] = 1;  
            // notify  
        } finally {  
            lock.unlock();  
        }  
    }  
}
```

```
class Consumer {  
    public void consume() {  
        try {  
            lock.lock();  
            while (isEmpty(buffer))  
                // wait  
            buffer[--count] = 0;  
            // notify  
        } finally {  
            lock.unlock();  
        }  
    }  
}
```



```
Lock lock = new ReentrantLock();
```

```
class Producer {  
    public void produce() {  
        try {  
            lock.lock();  
            while (isFull(buffer))  
                // wait  
            buffer[count++] = 1;  
            // notify  
        } finally {  
            lock.unlock();  
        }  
    }  
}
```

```
class Consumer {  
    public void consume() {  
        try {  
            lock.lock();  
            while (isEmpty(buffer))  
                // wait  
            buffer[--count] = 0;  
            // notify  
        } finally {  
            lock.unlock();  
        }  
    }  
}
```



```
Lock lock = new ReentrantLock();  
Condition notFull = lock.newCondition();
```

```
class Producer {  
    public void produce() {  
        try {  
            lock.lock();  
            while (isFull(buffer))  
                notFull.await();  
            buffer[count++] = 1;  
            // notify  
        } finally {  
            lock.unlock();  
        }  
    }  
}
```

```
class Consumer {  
    public void consume() {  
        try {  
            lock.lock();  
            while (isEmpty(buffer))  
                // wait  
            buffer[--count] = 0;  
            notFull.signal();  
        } finally {  
            lock.unlock();  
        }  
    }  
}
```



```
Lock lock = new ReentrantLock();  
Condition notFull = lock.newCondition();  
Condition notEmpty = lock.newCondition();
```

```
class Producer {  
    public void produce() {  
        try {  
            lock.lock();  
            while (isFull(buffer))  
                notFull.await();  
            buffer[count++] = 1;  
            notEmpty.signal();  
        } finally {  
            lock.unlock();  
        }  
    }  
}
```

```
class Consumer {  
    public void consume() {  
        try {  
            lock.lock();  
            while (isEmpty(buffer))  
                notEmpty.await();  
            buffer[--count] = 0;  
            notFull.signal();  
        } finally {  
            lock.unlock();  
        }  
    }  
}
```





```
Lock lock = new ReentrantLock();
Condition notFull = lock.newCondition();
Condition notEmpty = lock.newCondition();
```

```
class Producer {
    public void produce() {
        try {
            lock.lock();
            while (isFull(buffer))
                notFull.await();
            buffer[count++] = 1;
            notEmpty.signal();
        } finally {
            lock.unlock();
        }
    }
}
```

```
class Consumer {
    public void consume() {
        try {
            lock.lock();
            while (isEmpty(buffer))
                notEmpty.await();
            buffer[--count] = 0;
            notFull.signal();
        } finally {
            lock.unlock();
        }
    }
}
```




# Condition

A Condition object is used to park and awake threads

A Lock object can have any number of Conditions

A Condition object extends Object, so it has wait() and notify() methods... not to be taken for await() and signal()!





The `await()` call is blocking, can be interrupted

There are five versions for `await`:

- `await()`
- `await(time, TimeUnit)`
- `awaitNanos(nanosTimeout)`
- `awaitUntil(date)`
- `awaitUninterruptibly()`

These are ways to prevent the blocking of waiting threads with the `Condition` API

A fair Lock generates fair Condition

# Lock & Condition

Another implementation of the wait / notify pattern

It gives wiggle room to build better concurrent systems

Controlled interruptability, time out, fairness



# Read / Write Locks

---





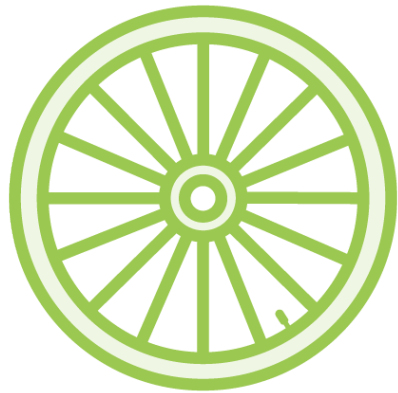
In some cases what we need is exclusive writes

And allow parallel reads

This is not how regular locks work

This is what read / write lock does






ReadWriteLock is an interface with only two methods

readLock() to get a read lock

writeLock() to get a write lock

Both are instances of Lock





Only one thread can hold the **write** lock

When the write lock is held, no one can hold the **read** lock

As many threads as needed can hold the **read** lock





```
ReadWriteLock readWritelock = new ReentrantReadWriteLock();
```

```
Lock readLock = readWritelock.readLock();
```

```
Lock writeLock = readWritelock.writeLock();
```

The right pattern is to create a read / write lock

Then get a read and a write lock as a pair from this read / write lock



```
Map<Long, User> cache = new HashMap<>();
```

```
try {  
    readLock.lock();  
    return cache.get(key);  
} finally {  
    readLock.unlock();  
}
```

It can be used to create a thread safe cache

Read-locking the get operation allows for concurrent reads



```
Map<Long, User> cache = new HashMap<>();
```

```
try {  
    writeLock.lock();  
    cache.put(key, value);  
} finally {  
    writeLock.unlock();  
}
```

It can be used to create a thread safe cache

Write-locking the put operation protects the non-concurrent map against internal corruption

It can also be achieved with a ConcurrentHashMap



# Read / Write Locks

Works with a single read / write lock object used to get a write lock and a read lock

Write operations are exclusive of other writes and reads

Read operations can be made in parallel


It thus allows for superior throughput



# Semaphores

---





A Semaphore is a well-known concept in concurrent programming

It looks like a lock, but allows several threads in the same block of code



```
Semaphore semaphore = new Semaphore(5); // permits
try {
    semaphore.acquire();
    // guarded block of code
} finally {
    semaphore.release();
}
```

A Semaphore built in the normal way is non-fair

The acquire() call is blocking until a permit is available

At most 5 threads can execute the guarded code at the same time



```
Semaphore semaphore = new Semaphore(5, true); // fair
try {
    semaphore.acquire(2);
    // guarded block of code
} finally {
    semaphore.release(2);
}
```

A Semaphore can be fair

The acquire() can ask for more than one permit

Then the release() call must release them all





```
Semaphore semaphore = new Semaphore(5);  
try {  
    semaphore.acquireUninterruptibly();  
    // guarded block of code  
} finally {  
    semaphore.release();  
}
```

By default if a waiting thread is interrupted it will throw an InterruptedException

Uninterruptibility means that the thread cannot be interrupted

It can be only be freed by calling its release() method



```
Semaphore semaphore = new Semaphore(5);  
try {  
    if (semaphore.tryAcquire())  
        // guarded block of code  
    else  
        // I could not enter the guarded code  
} finally {  
    semaphore.release();  
}
```

Acquisition can be made immediate...

... and can fail if there is already a thread in the guarded block of code



```
Semaphore semaphore = new Semaphore(5);
try {
    if (semaphore.tryAcquire(1, TimeUnit.SECONDS))
        // guarded block of code
    else
        // I could not enter the guarded code
} finally {
    semaphore.release();
}
```

One can also set a timeout before failing to acquire a permit

This pattern can also request more than one permit



# Handling Permits and Waiting Threads

One can **reduce** the number of permits (cannot increase it)

One can check the **waiting** threads:

- are there **any** waiting threads?
- **how many** threads are waiting?
- get the **collection** of the waiting threads



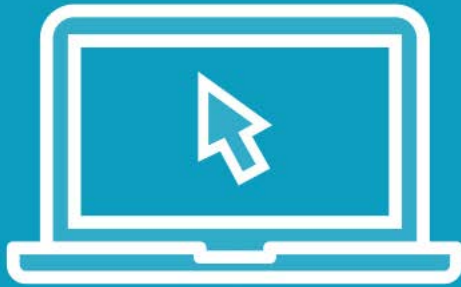
# Semaphore

It has a number of permits, which can be acquired in different ways, and released

One can query a semaphore for the number of waiting threads



# Demo



Let us see some code!

Let us see this producer / consumer pattern in action!

And how read / write lock can be used to create a concurrent cache



# Demo Wrapup



What did we see?

How to properly lock a Producer / Consumer with Lock

How to deal with exceptions and timeout

How HashMap can fail in a concurrent application

How to properly synchronize it using a pair of read / write locks



# Module Wrapup



What did we learn?

The difference between intrinsic and explicit locking

Give wiggle room to create efficient concurrent applications

Locks, read / write locks

Semaphores

Interruptability, time outs, fairness

